



Index Construction

Younghoon Kim

(nongaussian@gmail.com)

It contains some modified slides, which were originally written by Jim Martin, Donald Patterson Min-Yen Kan, and Zhang & Helmer, used for the Stanford CS276 class and from the Stuttgart IIR class

<https://nlp.stanford.edu/IR-book/newslides.html>



Index construction

- How do we construct an index?
- What strategies can we use **with limited main memory?**




Hardware basics

- Many design decisions in information retrieval are based on the characteristics of hardware
- We begin by reviewing hardware basics

DISK I/O

Storage Hierarchy



Storage Hierarchy (Cont.)


- **primary storage**: Fastest media but volatile (cache, main memory).
- **secondary storage**: next level in hierarchy, non-volatile, moderately fast access time
 - also called **on-line storage**
 - E.g. flash memory, magnetic disks
- ~~**tertiary storage**: lowest level in hierarchy, non-volatile, slow access time~~
 - ~~— also called **off-line storage**~~
 - ~~— E.g. magnetic tape, optical storage~~



Magnetic Hard Disk Mechanism




© CanStockPhoto.com - csp7646757





Magnetic Hard Disk Mechanism



NOTE: Diagram is schematic, and simplifies the structure of actual disk drives



Magnetic Disks

- **Read-write head**
 - Positioned very close to the platter surface (almost touching it)
 - Reads or writes magnetically encoded information.
- Surface of platter divided into circular **tracks**
 - Over 50K-100K tracks per platter on typical hard disks
- Each track is divided into **sectors**.
 - **A sector is the smallest unit of data that can be read or written.**
 - Sector size typically 512 bytes
 - Typical sectors per track: 500 to 1000 (on inner tracks) to 1000 to 2000 (on outer tracks)
- **To read/write a sector**
 - disk arm swings to position head on right track
 - platter spins continually; data is read/written as sector passes under head
- Head-disk assemblies
 - multiple disk platters on a single spindle (1 to 5 usually)
 - one head per platter, mounted on a common arm.



Performance Measures of Disks

- **Access time (=seek time)** – the time it takes from when a read or write request is issued to when data transfer begins. Consists of:
 - **Seek time** – time it takes to reposition the arm over the correct track.
 - Average seek time is 1/2 the worst case seek time.
 - 4 to 10 milliseconds on typical disks
 - **Rotational latency** – time it takes for the sector to be accessed to appear under the head.
 - Average latency is 1/2 of the worst case latency.
 - 4 to 11 milliseconds on typical disks (5400 to 15000 r.p.m.)
- **Data-transfer rate** – the rate at which data can be retrieved from or stored to the disk.
 - 25 to 100 MB per second max rate, lower for inner tracks
 - Multiple disks may share a controller, so rate that controller can handle is also important
 - E.g. SATA: 150 MB/sec, SATA-II 3Gb (300 MB/sec)
 - Ultra 320 SCSI: 320 MB/s, SAS (3 to 6 Gb/sec)



Optimization of Disk-Block Access

- **Block** – a contiguous sequence of sectors from a single track
 - data is transferred between disk and main memory in blocks
 - sizes range from 512 bytes to several kilobytes
 - Smaller blocks: more transfers from disk
 - Larger blocks: more space wasted due to partially filled blocks
 - Typical block sizes today range from 4 to 16 kilobytes



Optimization of Disk Block Access

- **File organization** – optimize block access time by organizing the blocks to correspond to how data will be accessed
 - E.g. Store related information on the same or nearby cylinders.
 - Files may get fragmented over time
 - E.g., if data is inserted to/deleted from the file
 - Or free blocks on disk are scattered, and newly created file has its blocks scattered over the disk
 - Sequential access to a fragmented file results in increased disk arm movement
 - Some systems have utilities to defragment the file system, in order to speed up file access



Optimization of Disk Block Access

head



s $+b$




$+s$ $+2b$



$+s$ $+b$



$+s$ $+b$




$4s+5b$

head



s $+5b$



$s+5b$

SORT-BASED INDEX CONSTRUCTION



Index Construction

- Grouping postings with term
 - → sort
- Scaling sort for large data
 - → external merge-sort

Scaling Index Construction

- Grouping postings

```
dic = dict()

for token, docid in postings:
    if token not in dic:
        dic[token] = list()

    dic[token].append(docid)
```

Term	docID
I	1
did	1
enact	1
julius	1


brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambit	2



Term	docID
ambit	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
I	1
it	2
julius	1
kill	1
kill	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2



Scaling Index Construction





Scaling Index Construction


- In-memory index construction does not scale
 - Can't stuff entire collection into memory, sort, then write back
- **How can we construct an index for very large collections?**
 - Considering the hardware constraints we just learned about . . .
 - Memory, disk, speed, etc.

Let's talk about straightforward methods

Sort using disk as “memory”?

- Can we use the same index construction algorithm for larger collections, but by using disk instead of memory?

head



- No: Sorting on disk is too slow due to **too many disk seeks**.




Sort using disk as “memory”?

- Block-based disk I/O



head



BLOCK-BASED DISK I/O



Example: Join Operation

- Join
 - Two relations 'student' and 'takes' on student IDs
 - E.g., `SELECT r.id FROM student as r JOIN takes as s ON r.id = s.studentID`
- Examples use the following information
 - Number of records of *student*. 5,000 *takes*. 10,000
 - Number of blocks of *student*. 100 *takes*. 400
 - That is,
 - A block can hold 50 student records and 25 takes records respectively
 - $n_r = 5,000$
 - $n_s = 10,000$
 - $b_r = 100$
 - $b_s = 400$



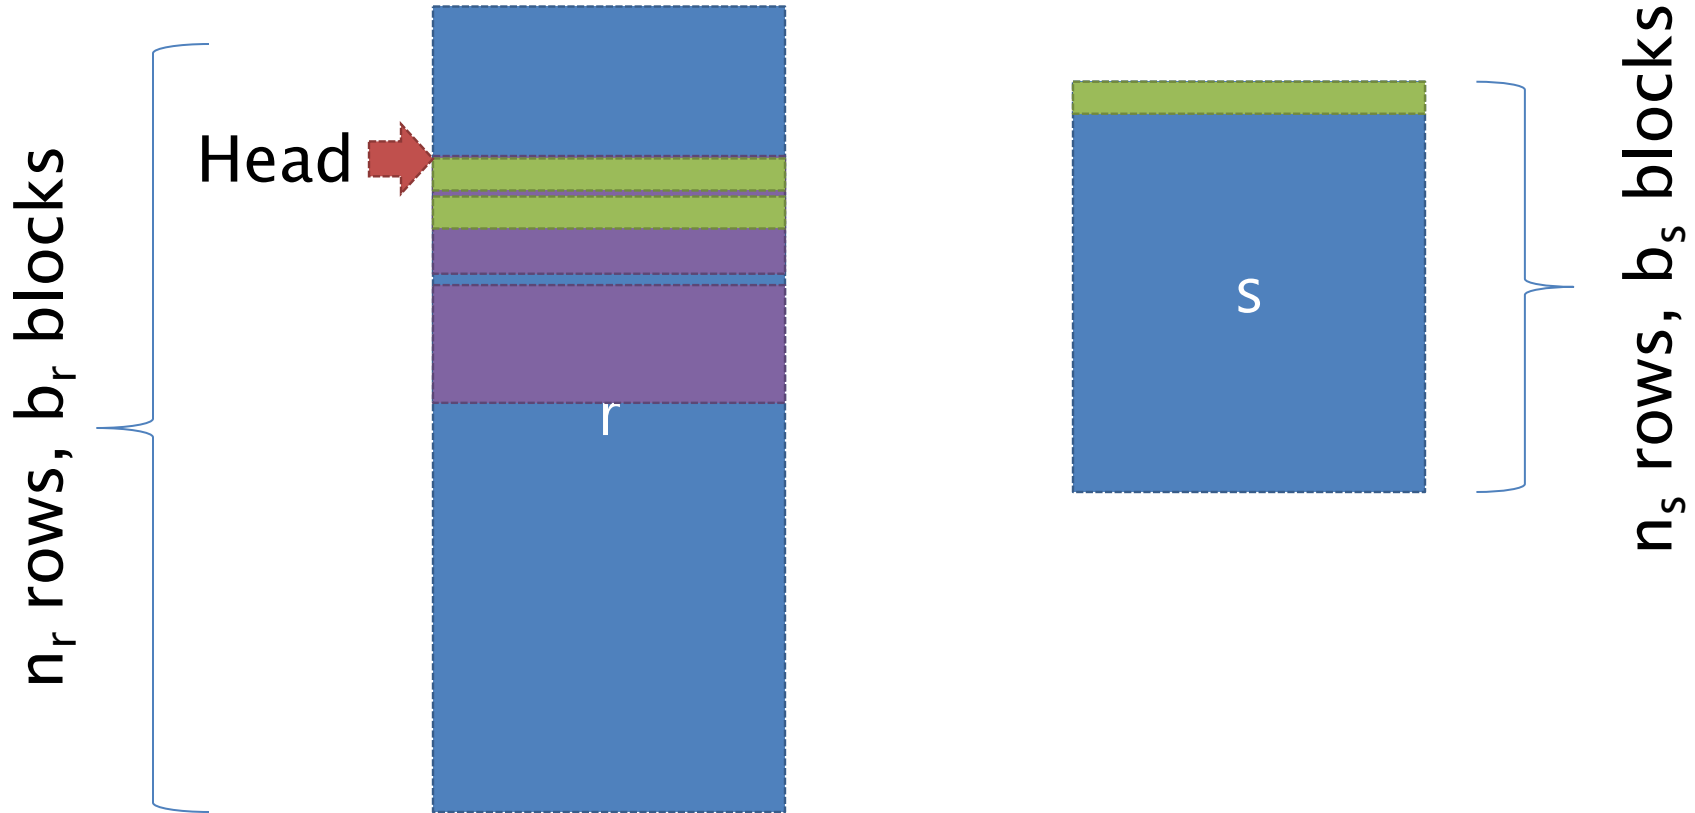
Nested-Loop Join

- To compute the theta join $r \bowtie_{\theta} s$ (θ : join condition)

```
for each tuple  $t_r$  in  $r$  do begin  
  for each tuple  $t_s$  in  $s$  do begin  
    test pair  $(t_r, t_s)$  to see if they satisfy  $\theta$   
    if they do, add  $(t_r, t_s)$  to the result  
  end  
end
```

- r is called the **outer relation** and s the **inner relation** of the join
- Assume that no indices can be used with any kind of join condition

Nested-Loop Join



of block transfers: $b_r + n_r * b_s$

of disk seeks: $b_r + n_r$



Nested-Loop Join (Cont.)

- In the worst case, if there is enough memory only to hold one block of each relation, the estimated cost is
$$\begin{array}{ll} n_r * b_s + b_r & \text{block transfers, plus} \\ n_r + b_r & \text{seeks} \end{array}$$
- Assuming worst case memory availability cost estimate is
 - with *student* as outer relation:
 - $5000 * 400 + 100 = 2,000,100$ block transfers,
 - $5000 + 100 = 5100$ seeks
 - with *takes* as the outer relation
 - $10000 * 100 + 400 = 1,000,400$ block transfers and 10,400 seeks
- Block nested-loops algorithm (next slide) is preferable.



Block Nested-Loop Join

- Variant of nested-loop join in which every block of inner relation is paired with every block of outer relation.

for each block B_r **of** r **do begin**

for each block B_s **of** s **do begin**

for each tuple t_r **in** B_r **do begin**

for each tuple t_s **in** B_s **do begin**

 Check if (t_r, t_s) satisfy the join condition

 if they do, add (t_r, t_s) to the result

end

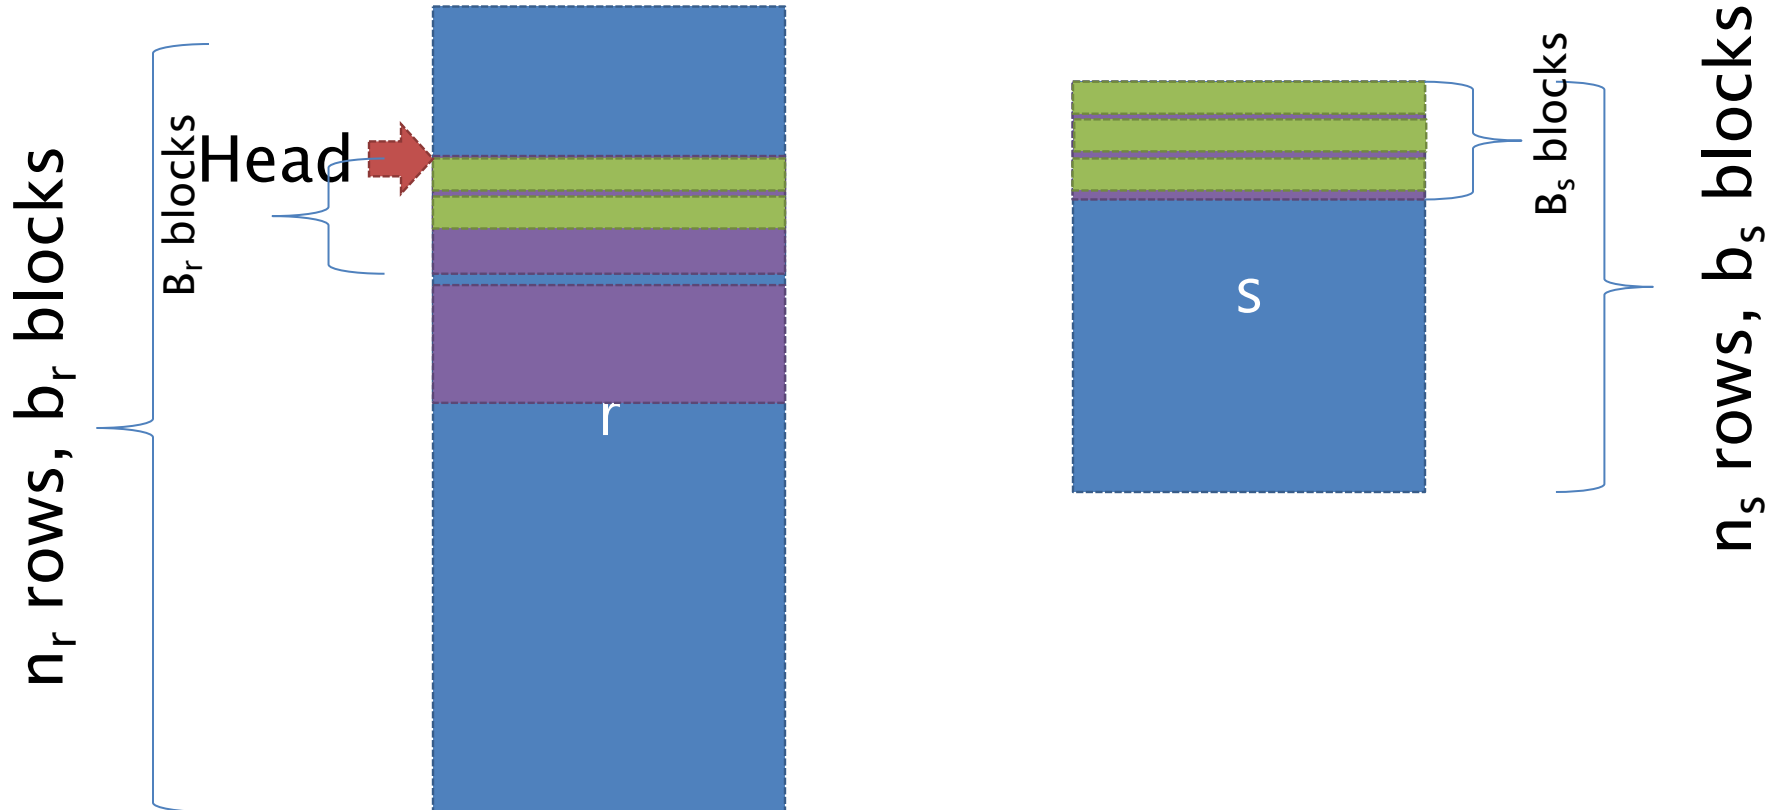
end

end

end

Block Nested-Loop Join

Available blocks: $B_r + B_s$



of block transfers: b_r + $b_s \cdot \left\lceil \frac{b_r}{B_r} \right\rceil$

of disk seeks: $\left\lceil \frac{b_r}{B_r} \right\rceil$ + $\left\lceil \frac{b_r}{B_r} \right\rceil$



Block Nested-Loop Join (Cont.)

- Available memory: $B_r + B_s$ blocks

$$\text{Block transfers: } b_r + b_s \cdot \left\lceil \frac{b_r}{B_r} \right\rceil$$

$$\text{Disk seeks: } 2 \left\lceil \frac{b_r}{B_r} \right\rceil$$

- In the worst case: $B_r = 1$ and $B_s = 1$
 - $b_r * b_s + b_r$ block transfers + $2 * b_r$ seeks
 - $100 * 400 + 100 = 40,100$ block transfers
 - $2 * 100 = 200$ seeks

Nested-loop Join:

2,000,100 block transfers +
5,100 seeks



Block Nested-Loop Join (Cont.)

- I/O costs

$$\text{Block transfers: } b_r + b_s \cdot \left\lceil \frac{b_r}{B_r} \right\rceil$$

$$\text{Disk seeks: } 2 \left\lceil \frac{b_r}{B_r} \right\rceil$$


- Improvements to nested loop and block nested loop algorithms:
 - In block nested-loop, use $(M - 2)$ memory blocks as the buffer for outer relations, where M = memory size in blocks; use remaining two blocks to buffer inner relation and output
 - **Cost = $b_r + b_s \cdot \left\lceil \frac{b_r}{M-2} \right\rceil$ block transfers + $2 \left\lceil \frac{b_r}{M-2} \right\rceil$ seeks**

EXTERNAL MERGE SORT

BSBI: Blocked sort-based Indexing

(Sorting with fewer *disk seeks*)


- 12-byte (4+4+4) records (*termID*, *docID*, *pos*).
- Run(Block aligned): Already sorted list of records
- Basic idea of algorithm:
 - Accumulate postings for each block, sort, write to disk.
 - Then merge the blocks into one long sorted order.



I/O Cost of Merging Two Runs

- Assume


- We merge two runs and write the intermediated sorted run on disk
- We are available $M (=2B + 1)$ blocks of memory



Sorting N blocks of records

- Mergesort do **binary merges**, with a merge tree of $\log_2 N$ depth where N = the number of smallest runs
- During each level, read into memory runs, merge, write back to disk
- Assume $M (=2B + 1)$ blocks of memory available and the initial run size is R blocks

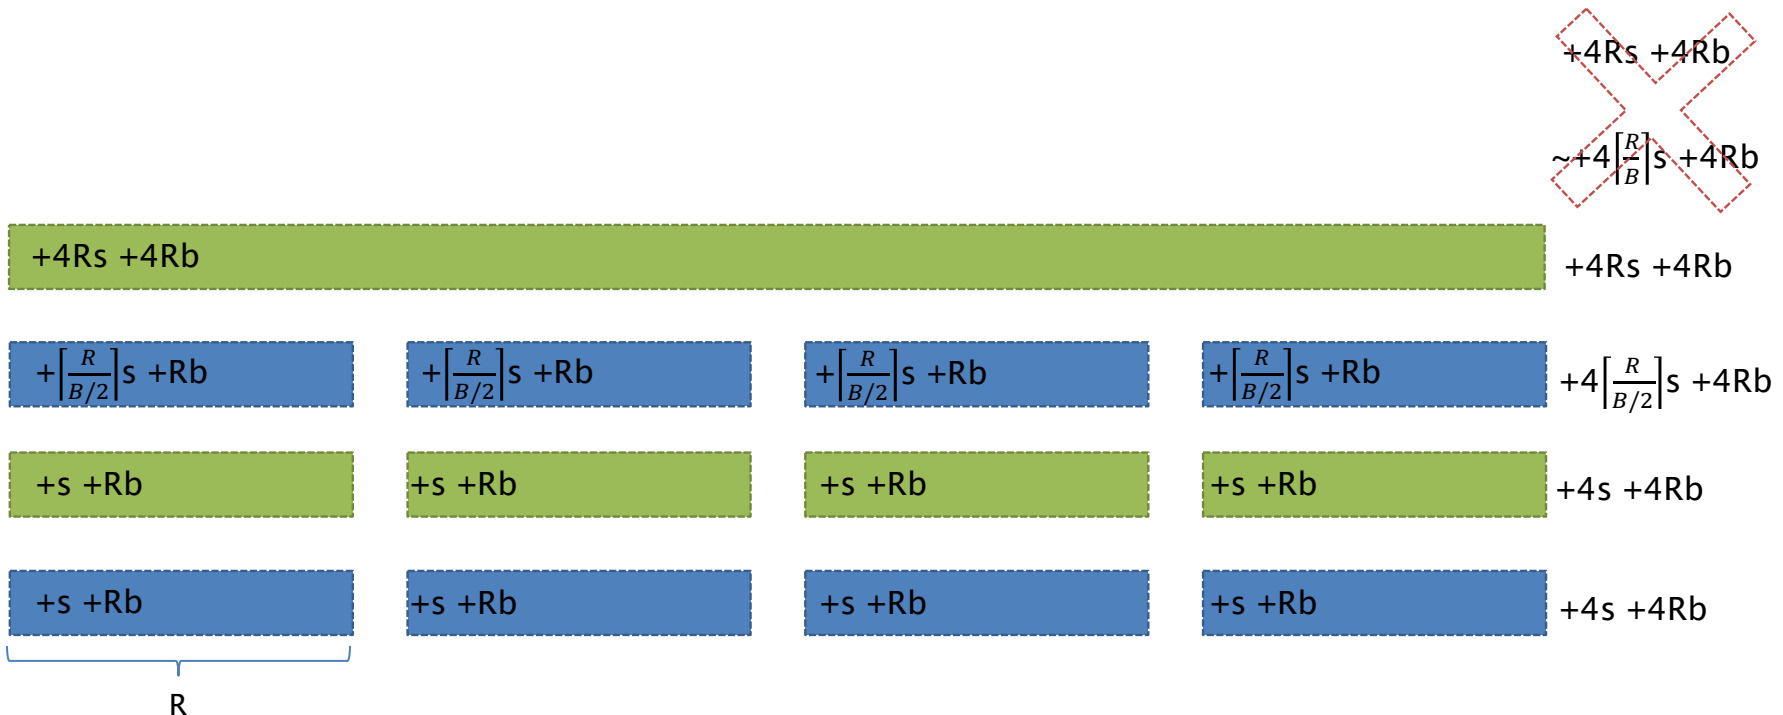
Temporary setting





Binary Merge vs. Multi-way Merge

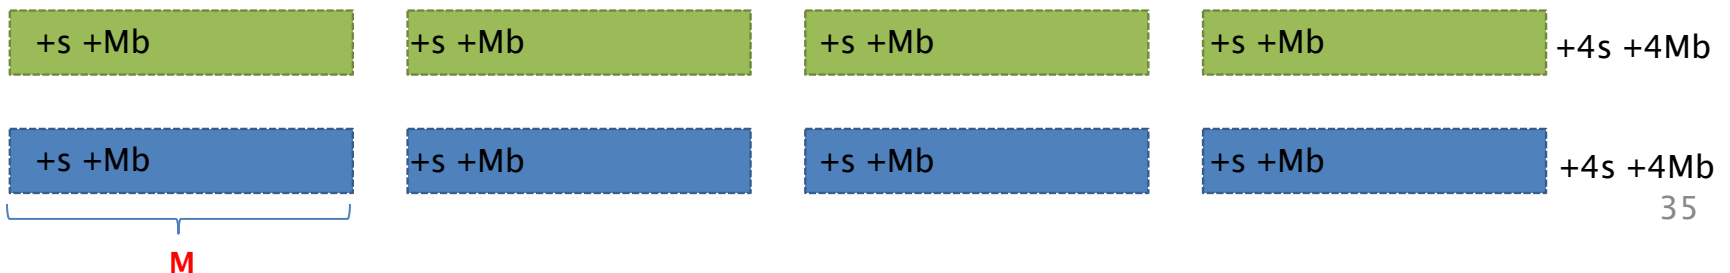
- But it is more efficient to do a multi-way merge, where you are reading from all blocks simultaneously
- Assume $M (=2B + 1)$ blocks of memory available and the initial run size is R blocks
 - Use $B/2$ blocks for each run






Multi-way Merge: Determine R

- Given
 - M blocks of available memory
 - m: m-way merge
- The depth of merge tree
 - $\log_m n_R$ where n_R is the number of initial runs
 - To reduce n_R , we use M blocks for sorting each initial run and generate M-blocks long initial runs



Multi-way Merge: Determine m

- Given
 - M blocks of available memory
 - m: m-way merge
- The depth of merge tree
 - $\log_m n_R$ where n_R is the number of initial runs
 - To increase m, we perform (M-1)-way merge: use 1-blocks for the read buffer for each run




Multi-way Merge

- Access m runs

- Use a block of memory for read buffer for each run
- Use a block of memory for output buffer
- Among the head of m runs, takes the smallest and put into output buffer



Output buffer: 1 1




Multi-way Merge

- Access m runs

- Use a block of memory for read buffer for each run
- Use a block of memory for output buffer
- Among the head of m runs, takes the smallest and put into output buffer

1 1

Output buffer: 2 2




Multi-way Merge

- Access m runs

- Use a block of memory for read buffer for each run
- Use a block of memory for output buffer
- Among the head of m runs, takes the smallest and put into output buffer

1 1 2 2

Output buffer: 2 2




Multi-way Merge

- Access m runs

- Use a block of memory for read buffer for each run
- Use a block of memory for output buffer
- Among the head of m runs, takes the smallest and put into output buffer

1 1 2 2 2 2

Output buffer: 3 4




Multi-way Merge

- Access m runs

- Use a block of memory for read buffer for each run
- Use a block of memory for output buffer
- Among the head of m runs, takes the smallest and put into output buffer

1 1 2 2 2 2 3 4

Output buffer: 4 5




Multi-way Merge

- Access m runs

- Use a block of memory for read buffer for each run
- Use a block of memory for output buffer
- Among the head of m runs, takes the smallest and put into output buffer

1 1 2 2 2 2 3 4 4 5

Output buffer: 5 6




Multi-way Merge

- Access m runs

- Use a block of memory for read buffer for each run
- Use a block of memory for output buffer
- Among the head of m runs, takes the smallest and put into output buffer

1 1 2 2 2 2 3 4 4 5 5 6

Output buffer: 6 6



Multi-way Merge

- Access m runs

- Use a block of memory for read buffer for each run
- Use a block of memory for output buffer
- Among the head of m runs, takes the smallest and put into output buffer

1 1 2 2 2 2 3 4 4 5 5 6 6 6


Output buffer: 7 7




Top Down vs. Bottom Up

- Traditional memory-based binary merge sort algorithm
 - Top down is easier to implement
- External m-way merge sort algorithm
 - Bottom up is better

<Top down>



<Bottom up>





External Merge-Sort

- Create sorted initial runs. Let i be 0 initially.
Repeatedly do the following till the end of the data:
 - (a) Read **M blocks** of postings into memory
 - (b) Sort the in-memory blocks using quick sort.
 - (c) Write sorted data to run R_i on disk; increment i .**Let the final value of i be N**
- Merge the runs (next slide).....

Let M denote **AVAILABLE**
Blocks in main memory



External Merge-Sort (Cont.)

2. Merge the runs (N-way merge). If $N < M$ ($N \leq M-1$),

1. Use N blocks of memory to buffer input runs, and 1 block to buffer output. Read the first block of each run into its buffer page
2. **repeat**
 1. Select the first record (in sort order) among all buffer pages
 2. Write the record to the output buffer. If the output buffer is full write it to disk.
 3. **If** the buffer page becomes empty **then** read the next block (if any) of the run into the buffer.
3. **until** all input buffer pages are empty:

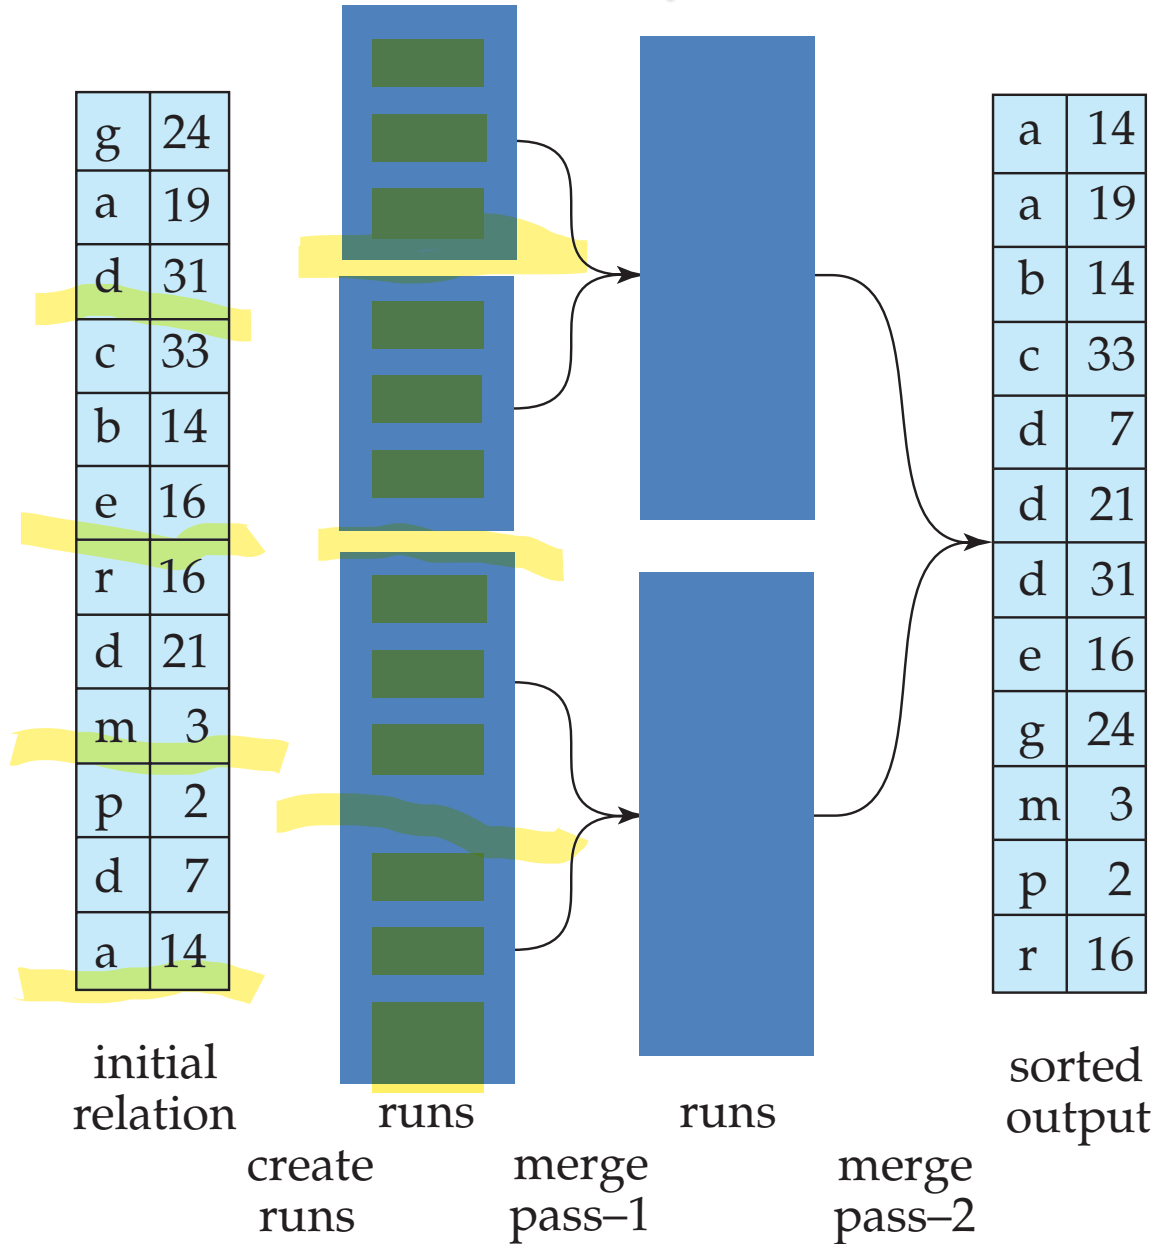


External Merge-Sort (Cont.)

- If $N \geq M$, **several merge passes** are required.
 - In each **pass**, contiguous groups of $M - 1$ runs are merged using the previous procedure.
 - A pass reduces the number of runs by a factor of $M - 1$ and creates runs longer by the same factor.
 - E.g. If there are 90 runs and $M=11$, one pass reduces the number of runs to 9 and each generated run is 10 times as long as the initial run
 - Repeated passes are performed till all runs have been merged into one.

Example:

M=3





Cost Estimation

■ Cost analysis:

- **M** blocks(=pages) are available in main memory
- Postings on b_r blocks $\rightarrow \left\lceil \frac{b_r}{M} \right\rceil$ runs are generated!
- Total number of **merge passes** required: $\left\lceil \log_{M-1} \left\lceil \frac{b_r}{M} \right\rceil \right\rceil$
- **COST**: Block transfers (read/write) for initial run creation as well as in each pass is $2b_r$
 - for final pass, we don't count write cost
 - ♦ we ignore final write cost for all operations since the output of an operation may be sent to the parent operation without being written to disk
 - Thus, total number of block transfers for external sorting:
$$2b_r + 2b_r \left(\left\lceil \log_{M-1} \left\lceil \frac{b_r}{M} \right\rceil \right\rceil \right) - b_r$$
- Seeks: next slide



Cost Estimation (Cont.)

■ Cost of seeks

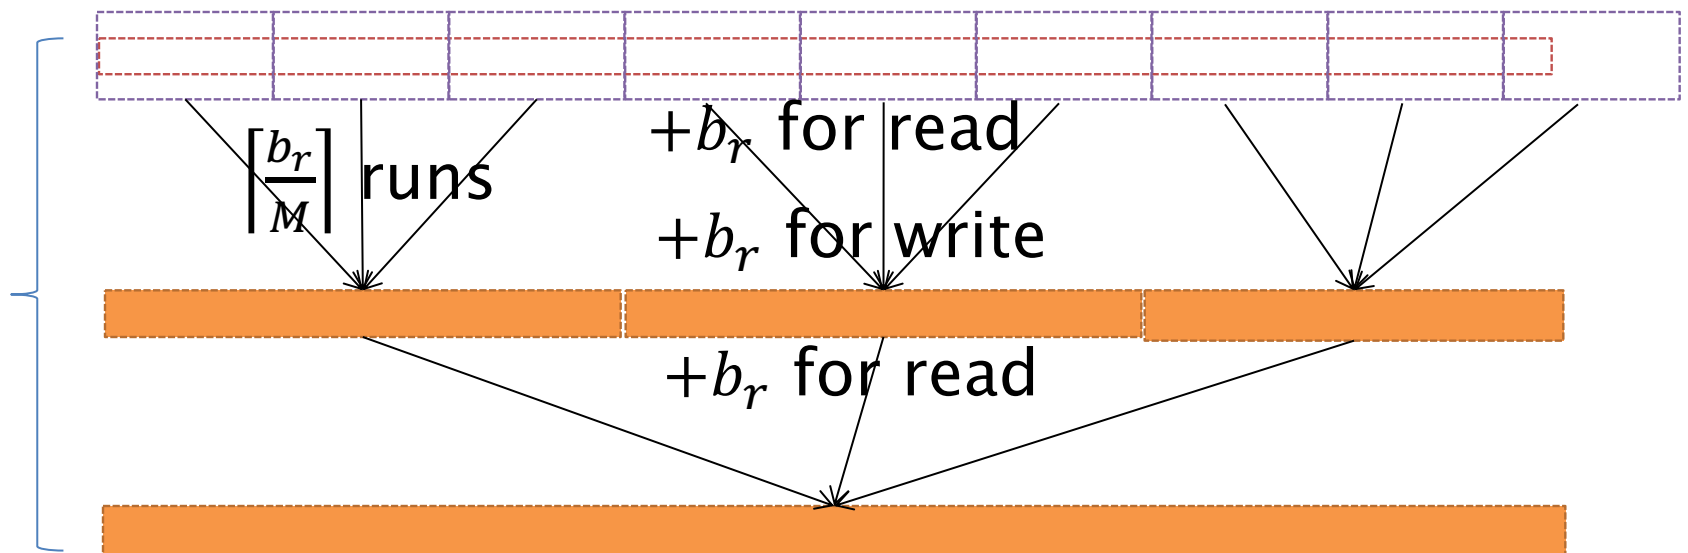
- **During run generation:** one seek to read each run and one seek to write each run
 - $2 \left\lceil \frac{b_r}{M} \right\rceil$ seeks
- **During the merge phase**
 - Need $2b_r$ seeks for each merge pass
 - ♦ except the final one which does not require a write
 - Total number of seeks:
$$2 \left\lceil \frac{b_r}{M} \right\rceil + 2b_r \left\lceil \log_{M-1} \left\lceil \frac{b_r}{M} \right\rceil \right\rceil - b_r$$

Cost Estimation: Block Transfers

- Cost analysis:

 M blocks

$2b_r$ for initial run generation



$\lceil \log_{M-1} \lceil \frac{b_r}{M} \rceil \rceil$ passes



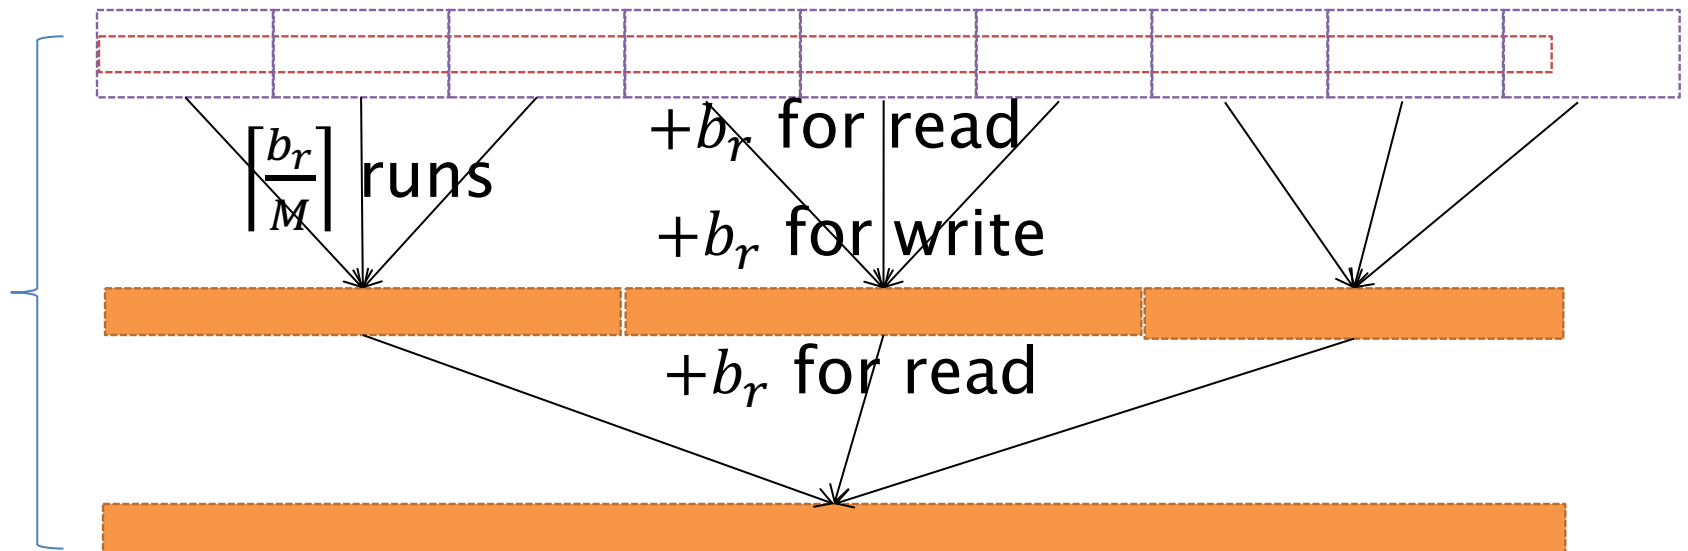
$$2b_r + 2b_r \left(\lceil \log_{M-1} \lceil \frac{b_r}{M} \rceil \rceil \right) - b_r$$

Cost Estimation: Disk Seeks

- Cost analysis:

 M blocks

$2 \left\lceil \frac{b_r}{M} \right\rceil$ for initial run generation




$\left\lceil \log_{M-1} \left\lceil \frac{b_r}{M} \right\rceil \right\rceil$ passes

$\Rightarrow 2 \left\lceil \frac{b_r}{M} \right\rceil + 2b_r \left(\left\lceil \log_{M-1} \left\lceil \frac{b_r}{M} \right\rceil \right\rceil \right) - b_r$

Exercise:

M=5



Exercise:

13	4
4	13
35	29
49	35
43	43
46	46
29	49
85	83
88	85
83	88
8	6
85	8
74	43
46	46
50	50
6	74
79	79
91	84
43	85
84	91
65	10
30	11
10	30
47	47
11	65
83	83

Output buffer



M=5 and 2 records per bloc



4way가 ㄴㅇ



1번만에 가ㄴㅇ

$$\rightarrow \frac{26}{10} = 3$$

Exercise

M=3 and 1 records per block

86
60
90
59
48
19
86
94
23
36
48
7
9
63
74
29
39
20

Quiz

- Sort the postings using external merge-sort algorithm
 - 30 postings
 - 3 blocks of memory are available
 - Each block(a page) can hold up to 3 postings
 - A block is used for the read buffer of each run
- **How many blocks to read/write?**
- **How many disk seeks occurs?**

Term	Doc #
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2
was	2



Distributed Indexing

- External merge-sort is the best choice using a single computer
- Nonetheless, web-scale indexing must use a distributed computing cluster
- How do we exploit such a pool of machines for indexing? → MapReduce