



# Parallel Indexing & MapReduce


Younghoon Kim

Hanyang University ERICA

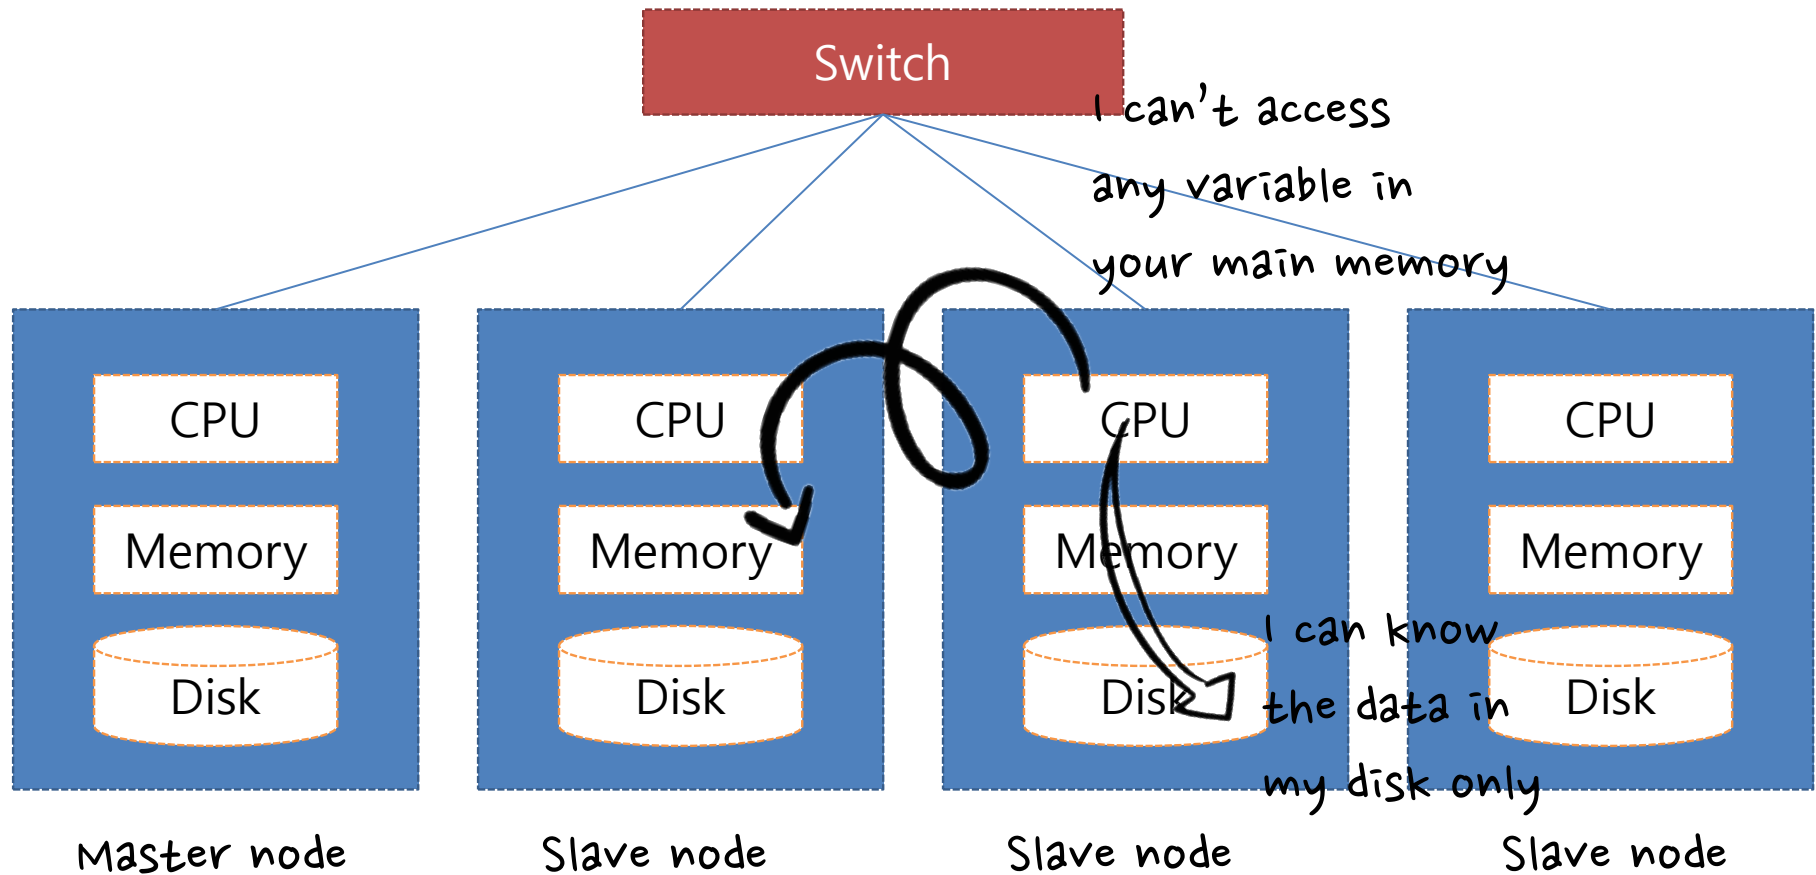
(nongaussian@hanyang.ac.kr)

# **PARALLEL PROGRAMMING USING MAPREDUCE**

# Single Node Architecture



# Shared-Nothing Cluster Architecture





# Programming Model

---

- Functional programming
- Users implement interface of two functions:
  - `map (in_key, in_value) ->`  
`(out_key, intermediate_value) list`
  - `reduce (out_key, intermediate_value list) ->`  
`out_value list`

# **MAP/REDUCE EXAMPLE #1**

## **(WORD COUNTING)**




# Word Counting

---

```
main () {  
    fd = open file ('big text file');  
    cnt = initialize a hash table;  
    while ( (line = read_a_line (fd)) != null) {  
        tokens = tokenize (line);  
        foreach (word in tokens) {  
            if (cnt[word] is defined) {  
                cnt[word] += 1;  
            }  
            else {  
                cnt[word] = 1;  
            }  
        }  
    }  
}
```

# Word Counting with MapReduce




- I can read only a line
- We cannot use any hash table
- What I can do is

```
tokens ← tokenize (line);  
foreach (word in tokens)  
{  
    emit(word, 1);  
}
```



# Word Counting with MapReduce



<apple, 1>  
<apple, 1>  
<apple, 1>  
<banana, 1>  
<banana, 1>  
<banana, 1>  
<banana, 1>  
<cat, 1>  
<cat, 1>  
...



# Word Counting with MapReduce

Scarabs  
소똥구리

apple cat banana

cat banana

apple banana

A very large text

apple banana

<apple, 1>

<apple, 1>

<apple, 1>

<banana, 1>

<banana, 1>

<banana, 1>

<banana, 1>

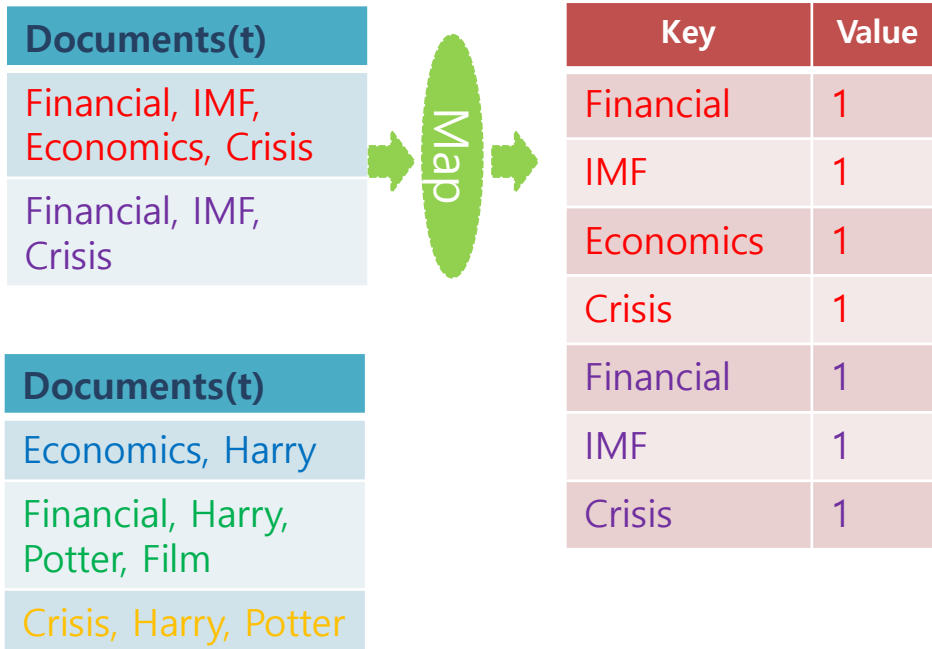
<cat, 1>

<cat, 1>

...



# Word Counting with MapReduce



# Word Counting with MapReduce

## Documents(t)

Financial, IMF,  
Economics, Crisis

Financial, IMF,  
Crisis

## Documents(t)

Economics, Harry

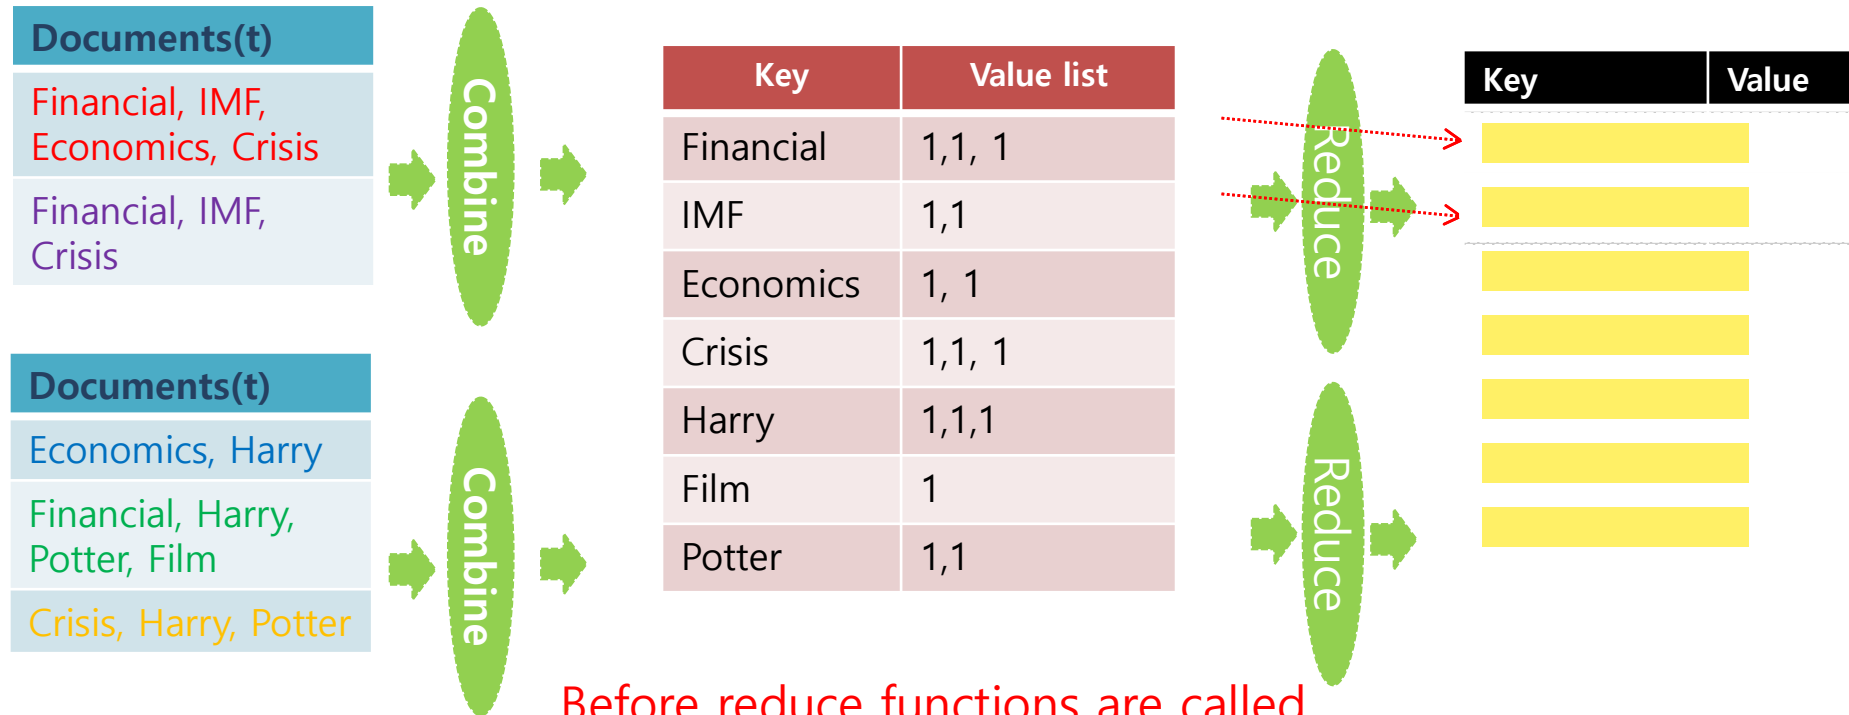
Financial, Harry,  
Potter, Film

Crisis, Harry, Potter



Key	Value
Economics	1
Harry	1
Financial	1
Harry	1
Potter	1
Film	1
Crisis	1
Harry	1
Potter	1

# Word Counting with MapReduce



Before reduce functions are called,  
for each distinct key,  
the list of its values are generated



# Hadoop MapReduce Programming in Java

```
public static class Map extends MapReduceBase implements
    Mapper<LongWritable, Text, Text, IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>
        output, Reporter reporter) throws IOException {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            output.collect(word, one);
        }
    }
}
```



# Hadoop MapReduce Programming in Java

```
public static class Reduce extends MapReduceBase implements
    Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text,
        IntWritable> output, Reporter reporter) throws IOException {
        int sum = 0;
        while (values.hasNext()) { sum += values.next().get(); }
        output.collect(key, new IntWritable(sum));
    }
}
```



# Hadoop MapReduce Programming in Java

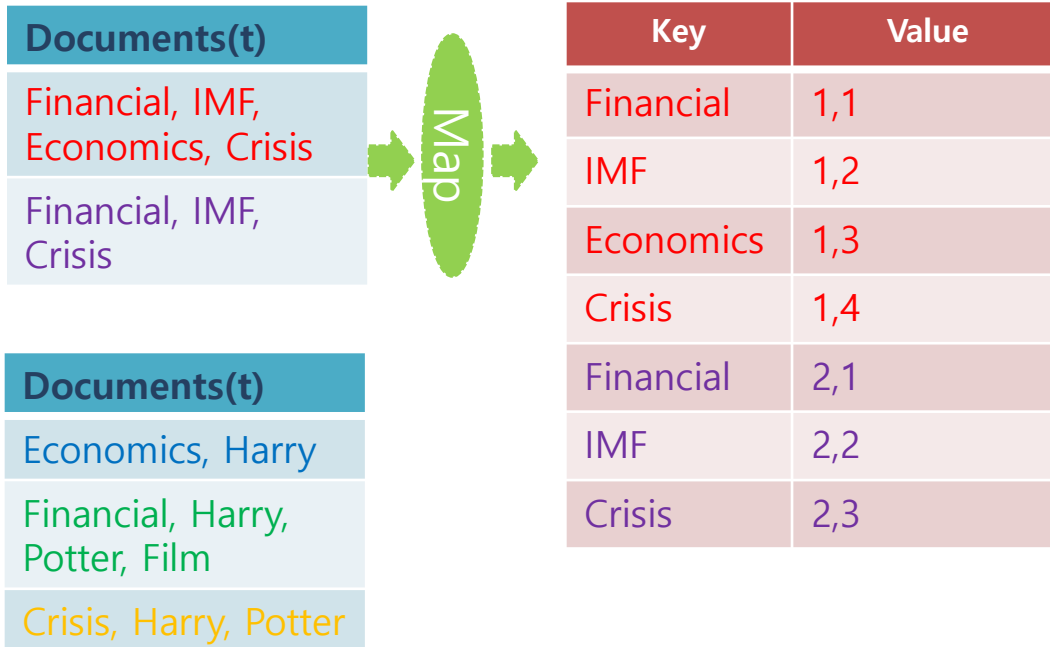
```
public static void main(String[] args) throws Exception {  
    JobConf conf = new JobConf(WordCount.class);  
    conf.setJobName("wordcount");  
    conf.setOutputKeyClass(Text.class);  
    conf.setOutputValueClass(IntWritable.class);  
    conf.setMapperClass(Map.class);  
    conf.setCombinerClass(Reduce.class);  
    conf.setReducerClass(Reduce.class);  
    conf.setInputFormat(TextInputFormat.class);  
    conf.setOutputFormat(TextOutputFormat.class);  
    FileInputFormat.setInputPaths(conf, new Path(args[0]));  
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));  
  
    JobClient.runJob(conf);  
}}
```



# **MAP/REDUCE EXAMPLE #2**

## **(BUILDING AN INVERTED INDEX)**

# An Example of Indexing





# An Example of Indexing

---

## Documents(t)

Financial, IMF,  
Economics, Crisis

Financial, IMF,  
Crisis

## Documents(t)

Economics, Harry

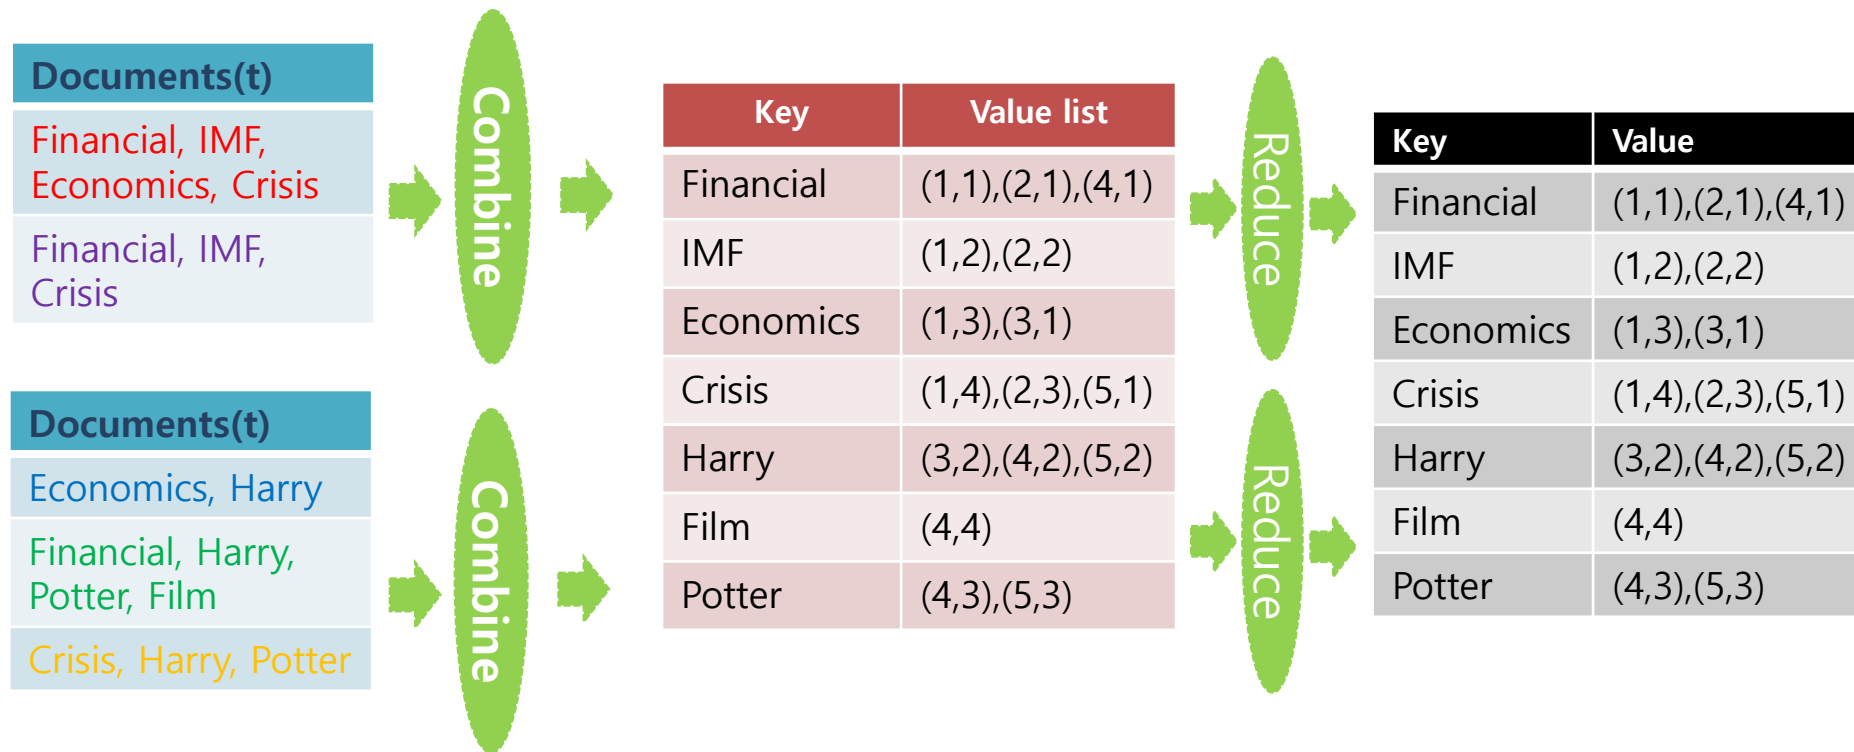
Financial, Harry,  
Potter, Film

Crisis, Harry, Potter



Key	Value
Economics	3,1
Harry	3,2
Financial	4,1
Harry	4,2
Potter	4,3
Film	4,4
Crisis	5,1
Harry	5,2
Potter	5,3


# An Example of Indexing



Before reduce functions are called,  
for each distinct key,  
the list of its values are generated

# **MAP/REDUCE EXAMPLE #3 (AGGREGATION IN NOSQL)**

Practice with



mongoDB



# Characteristics

Short for Binary JSON:  
Binary-encoded  
serialization

	mongoDB
Data Model	Document-oriented (based on BSON)
Interface	Custom protocol over TCP/IP
Object Storage	Database contains <b>collections (=tables)</b> Collections contains <b>documents (=rows)</b>
Query Method	<b>MapReduce (javascript)</b> creating collections + Object-cased query language
Replication	Master-Slave
Concurrency	Update in-place
Written In	C++



# Select-Where Query

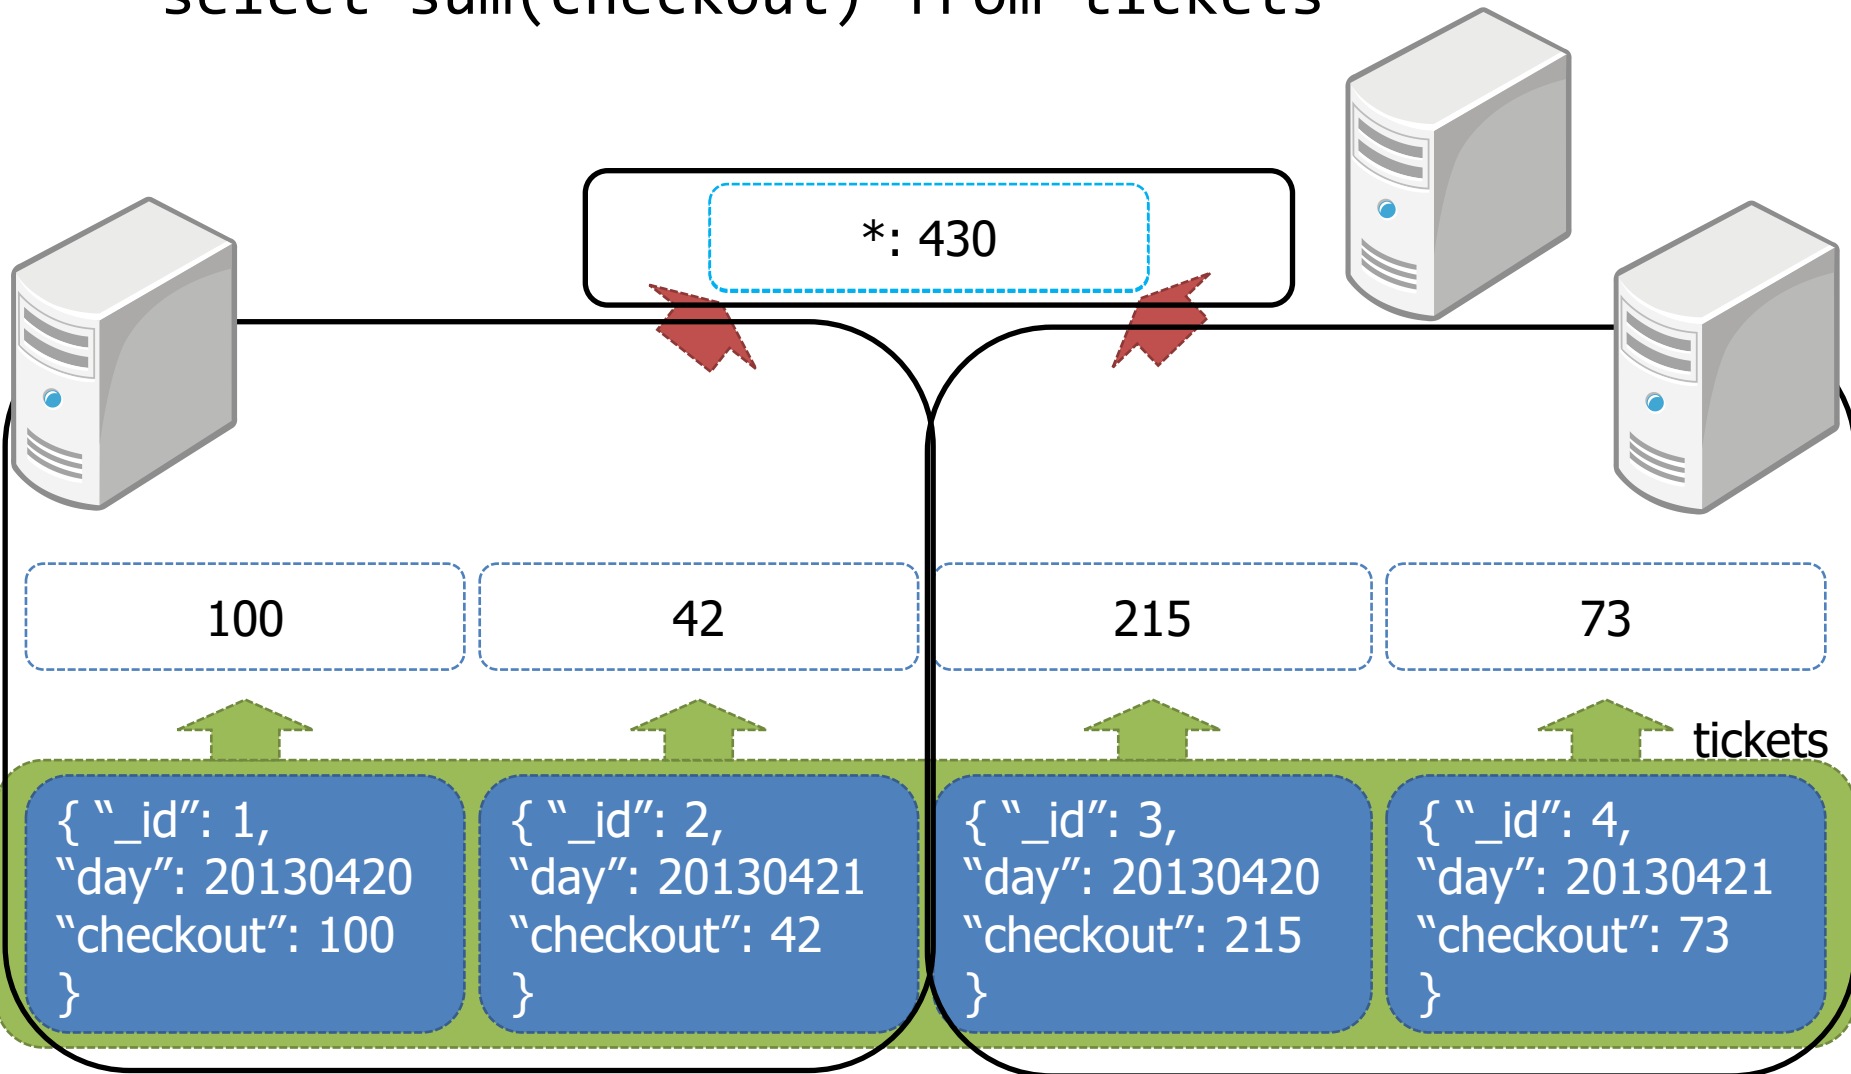
- Example:
  - select \* from colors where name='green';  
→ `db.colors.find({name:'green'})`
  - select \* from people where age <= 27;  
→ `db.people.find({age:{$lte:27}})`

연산자  
\$gt  
\$gte  
\$lt  
\$lte  
\$ne  
\$in  
...



# Summation with MapReduce

- `select sum(checkout) from tickets`





# Example

```
db.tickets.mapReduce(  
  function() {  
    emit( '*', this.checkout );  
  },  
  function(key, values) {  
    return Array.sum(values);  
  },  
  { out: 'postout', query: { status: 'active' } }  
) .find();
```

→ { "\_id" : "\*", "value" : 430 }



# MapReduce Commend

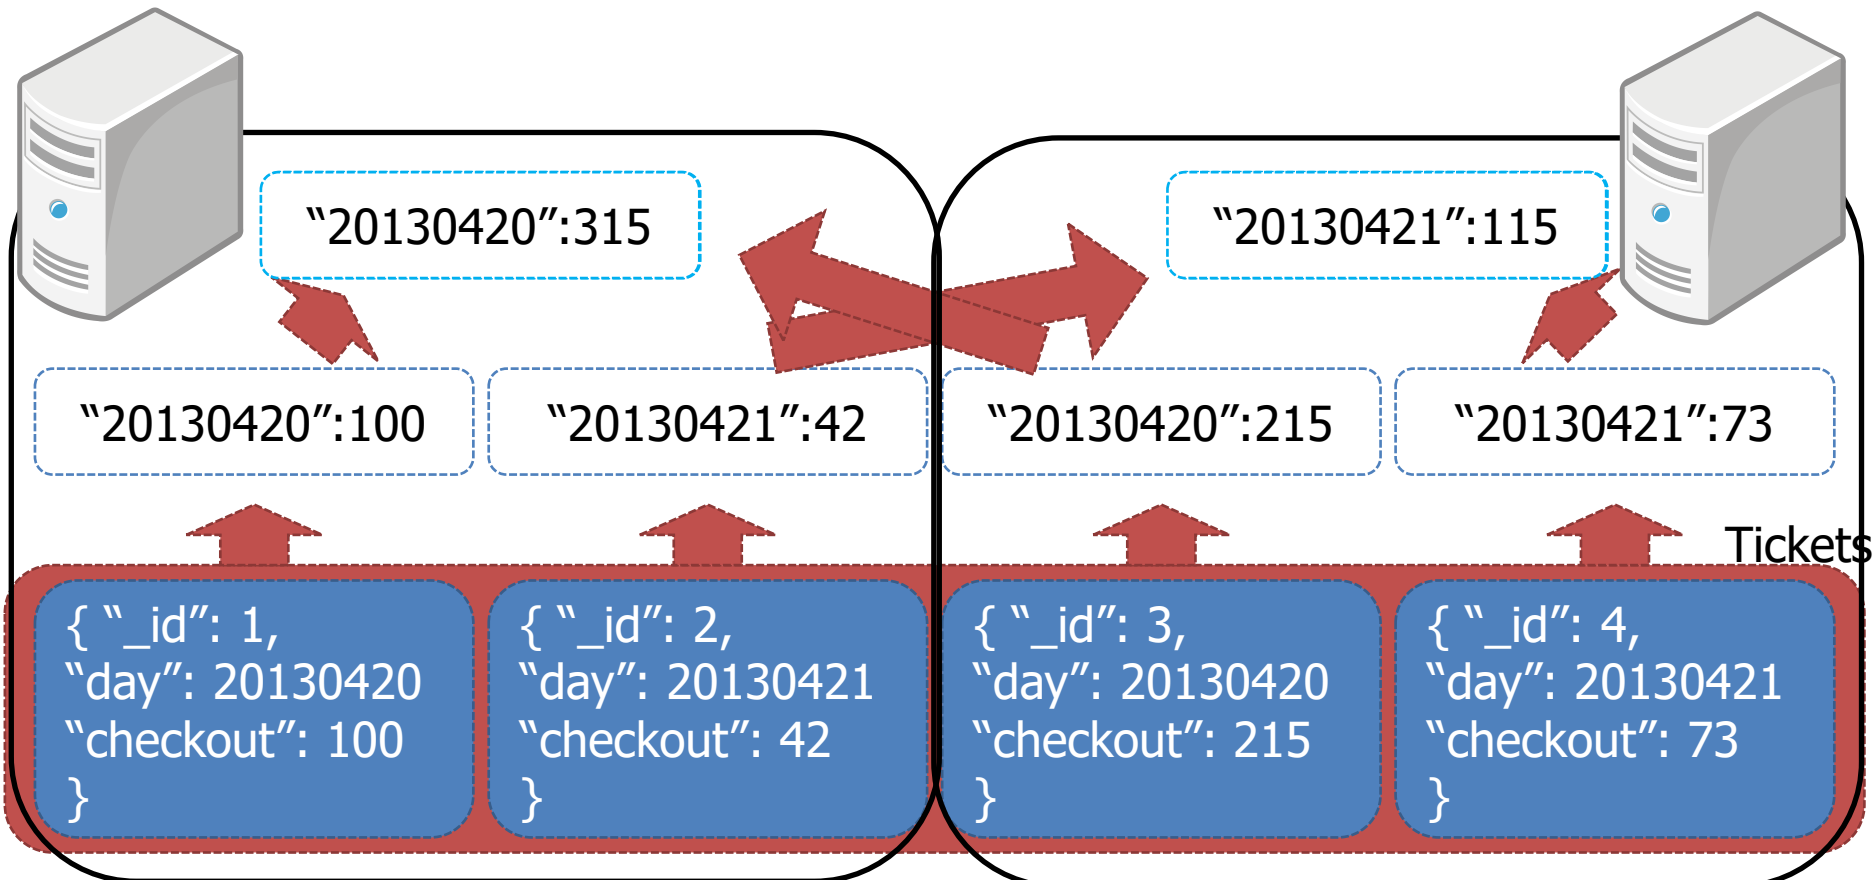
---

- `> db.collection.mapReduce(`
- `function() { ... emit(key,value); ... },`
- `function(key,values) {return output_value;},`
- `{`
  - `out: collection,`
  - `query: document,`
  - `sort: document,`
  - `limit: number`
- `})`

- `map` is a javascript function that maps a value with a key and emits a key-value pair
- `reduce` is a javascript function that reduces or groups all the documents having the same key
- `out` specifies the location of the map-reduce query result
- `query` specifies the optional selection criteria for selecting documents
- `sort` specifies the optional sort criteria
- `limit` specifies the optional maximum number of documents to be returned

# Groupby with MapReduce

- `select sum(checkout) from tickets group by day`





# Groupby with MapReduce

---

- ```
db.tickets.mapReduce(  
  function() {  
    emit( this.day, this.checkout);  
  },  
  function(key, values) {  
    return Array.sum(values);  
  },  
  {out:'groupby'}  
)
```

.find()
- { "\_id" : 20130420, "value" : 315 }
- { "\_id" : 20130421, "value" : 115 }

# **MAP/REDUCE EXAMPLE #4 (LOGISTIC REGRESSION WITH SPARK)**


Practice with



In-Memory cluster computing for  
Iterative and Interactive Applications

# Motivation

- Popular MapReduce implementations such as Hadoop transform data flowing from stable storage to stable storage







# Motivation

- Acyclic data flow is a powerful abstraction, but is not efficient for applications that repeatedly reuse a working set of data:
  - **Iterative algorithms** (many in machine learning, e.g., PageRank, EM algorithms)
  - **Interactive data mining tools** (R, Excel, Python)
- Spark introduces augment data flow model with “resilient distributed datasets” (RDDs)



# ~~Resilient distributed datasets (RDDs)~~

---

- An RDD is an immutable, partitioned, logical collection of records
  - Need not be materialized, but rather contains information to rebuild a dataset from stable storage
- Partitioning can be based on a key in each record (using hash or range partitioning)
- Built using bulk transformations on other RDDs
- Can be cached for future reuse



# Simple Spark Apps: Word Counting

Scala

RDD


```
val f = sc.textFile("README.md")
val wc = f.flatMap(l => l.split(" ")).map(word => (word, 1)).reduceByKey(_ + _)
wc.saveAsTextFile("wc_out")
```

Python:

```
from operator import add
f = sc.textFile("README.md")
wc = f.flatMap(lambda x: x.split(' ')).map(lambda x: (x, 1)).reduceByKey(add)
wc.saveAsTextFile("wc_out")
```


# Logistic Regression

- Goal: find the **best line separating** two sets of points



# Logistic Regression

- Goal: find the **best line separating** two sets of points





# Optimization Problem

- Maximize

$$\sum_{i=1}^n y_i \cdot \log\left(\frac{1}{1 + e^{-(\beta_0 + x_i \cdot \beta)}}\right) + \sum_{i=1}^n (1 - y_i) \cdot \log\left(1 - \frac{1}{1 + e^{-(\beta_0 + x_i \cdot \beta)}}\right)$$

- Gradient descent method

$$\beta^{(t+1)} \leftarrow \beta^{(t)} + \alpha \sum_{i=1}^n \left( y_i - \frac{1}{1 + e^{-(\beta_0 + x_i \cdot \beta)}} \right) x_i$$


Sum of values  
calculated with  
each data points



# Logistic Regression Code

- `val data = spark.textFile(...).map(readPoint).cache()`
- `var w = Vector.random(D)`
- `for (i <- 1 to ITERATIONS) {`
- `val gradient = data.map(p =>`
- `(p.y - 1 / (1 + exp(-(w dot p.x)))) * p.x`
- `).reduce(_ + _)`
- `w += gradient`
- `}`
- `println("Final w: " + w)`

# Logistic Regression Performance







# Summary

Programming with MapReduce is  
not a choice, but a necessity.  
Don't worry. It is fun!