
Real-Time Systems

C. L. Liu and J. W. Layland, "Scheduling Algorithms for Multiprogramming in a Hard Real-Time Environment", Journal of the ACM, 1973

Contents

- Background and Assumptions
- Terminologies: Revisited
- Rate-monotonic Priority Assignment
- Least Upper Bound (LUB) of Processor Utilization
- Some Results on Rate Monotonic Algorithm

Assumptions

- (A1) The requests for all tasks for which hard deadlines exist are periodic.
 - Only for periodic task set
- (A2) Deadlines consist of run-ability constraints only.
 - Each task must be completed before its next invocation.
- (A3) The tasks are independent to each other.
 - No precedence relation
- (A4) Run-time for each task is constant for that task and does not vary with time.
 - Assumes worst-case execution time known a priori.
- (A5) Any non-periodic tasks in the system are special.
 - Soft deadline; background jobs

Rate Monotonic RT Scheduling

- Basic rate monotonic scheduling (RMS)
 - Rate monotonic priority
 - The higher rate, the higher priority
 - Schedulability guaranteed if utilization rate is below a certain limit
 - For feasible schedules
 - $f_i = 1/T_i$: frequency (=rate)
 - C_i : execution time

$$\sum_{i=1}^n C_i f_i \leq 1$$

Rate Monotonic RT Scheduling

- If the total utilization rate has least upper bound (LUB) of $n(2^{1/n} - 1)$ where $n = \text{\#tasks}$, there exists a feasible rate monotonic schedule. That is,

$$\sum_{i=1}^n c_i f_i \leq n(2^{\frac{1}{n}} - 1) = U(n)$$

- $U(1) = 1.0$, $U(2) = 0.828$, $U(3) = 0.779$, $U(\infty) = \ln 2 = 0.693$
- Only sufficient condition
- Priority inversion problem

Properties of RM Scheduling

- On-line
- Preemptive
- Priority-driven
- Static (Fixed-priority)

Terminologies: Revisted

- Task set $\tau = \{ \tau_1, \tau_2, \dots, \tau_n \}$
 - A task is characterized by (T_i, C_i)
- Deadline of a request
 - The time of the next request for the same task
- An *overflow* occurs at time t if t is the deadline of an unfulfilled request
- For a given set of tasks, a scheduling algorithm is *feasible* if the tasks are scheduled so that no overflow ever occurs.
- Response time of a request for a certain task
 - The time span between the request and the end of the response to that request


Critical Instant

Definition

A ***critical instant*** for a task is an instant at which a request for that task will have the largest response time.

Theorem 1.

A critical instant for any task occurs whenever the task is requested simultaneously with requests for all higher priority tasks.



Priority Assignment

- From the proof of Theorem 1,
 - A simple direct calculation can determine whether or not a given priority assignment will yield a feasible scheduling algorithm.
 - If the request for all tasks at their critical instants are fulfilled before their deadlines, then the scheduling algorithm is feasible.

Priority Assignment (cont'd)


- “Reasonable” rule of priority assignment
 - Assign priorities to tasks according to their request rates, independent of their execution times.
- Rate monotonic (RM) priority assignment
 - Tasks with higher request rates (shorter period) will have higher priorities.

Theorem 2. Optimality of RM Priority Assignment

If a feasible priority assignment exists for some task set, the rate-monotonic priority is feasible for that task set.

Scheduling Example

- Three tasks $\tau_1=(2,0.5)$, $\tau_2=(6,2)$, $\tau_3=(10,1.75)$



LUB of Processor Utilization

- Utilization factor U
 - The fraction of processor time spent in the execution of the task set (= 1-(the fraction of idle processor time))
- Full utilization of processor
 - A set of tasks is said to “*fully utilize*” the processor if the priority assignment is feasible for the set and if an increase in the run-time of any of the tasks in the set will make the priority assignment infeasible
- LUB (least upper bound) of utilization factor
 - The minimum of utilization factor over all sets of tasks that fully utilize the processor
 - For all task sets whose utilization factor is below this bound, there exists a fixed priority assignment which is feasible

A schedule is said to be *feasible* if all tasks can be completed according to the set of specified constraints

A set of tasks is said to be *schedulable* if there exists at least one algorithm that can produce a feasible schedule

Assumptions of LUB


- The processor always executes the highest-priority task.
- Task priorities are assigned according to rate monotonic policy.
- Tasks do not synchronize with each other.
- Each task's deadline is at the end of its period.
- Tasks do not suspend themselves in the middle of computations.
- No interrupts, No context switching overhead

LUB of Processor Utilization (cont'd)

Theorem 3. LUB for a set of two tasks

For a set of two tasks with fixed priority assignment, the least upper bound to the processor utilization factor is $U=2(2^{1/2}-1)$.

Proof: (a) When $C_1 < T_2 - \lfloor T_2/T_1 \rfloor T_1$




$$C_2 = T_2 - C_1 \left\lceil \frac{T_2}{T_1} \right\rceil$$

$$U = \frac{C_1}{T_1} + \frac{C_2}{T_2} = \frac{C_1}{T_1} + 1 - \frac{C_1}{T_2} \left\lceil \frac{T_2}{T_1} \right\rceil$$

Since
$$C_1 \left(\frac{1}{T_1} - \frac{1}{T_2} \left\lceil \frac{T_2}{T_1} \right\rceil \right) \leq 0$$

The processor utilization is monotonically decreasing in C_1

LUB of Processor Utilization (cont'd)



$$C_2 = (T_1 - C_1) \left\lfloor \frac{T_2}{T_1} \right\rfloor$$

$$\begin{aligned} U &= \frac{C_1}{T_1} + \frac{C_2}{T_2} \\ &= \frac{C_1}{T_1} + \left\lfloor \frac{T_2}{T_1} \right\rfloor \left(\frac{T_1 - C_1}{T_2} \right) \\ &= \frac{T_1}{T_2} \left\lfloor \frac{T_2}{T_1} \right\rfloor + C_1 \left(\frac{1}{T_1} - \frac{1}{T_2} \left\lfloor \frac{T_2}{T_1} \right\rfloor \right) \end{aligned}$$

Since $\frac{1}{T_1} - \frac{1}{T_2} \left\lfloor \frac{T_2}{T_1} \right\rfloor \geq 0$

The processor utilization is monotonically increasing in C_1

LUB of Processor Utilization (cont'd)

Proof (cont'd):

In case (a), U is monotonically decreasing in C_1 .

In case (b), U is monotonically increasing in C_1 .

\therefore The minimum value of U occurs at $C_1 = T_2 - T_1 \lfloor T_2/T_1 \rfloor$.

$$\text{Let } I = \lfloor T_2/T_1 \rfloor \text{ and } T_2/T_1 = I + f (0 \leq f < 1) \Leftrightarrow \left\lceil \frac{T_2}{T_1} \right\rceil = \begin{cases} I & \text{if } f = 0 \\ I + 1 & \text{otherwise} \end{cases}$$

$$U = \frac{C_1}{T_1} + 1 - \frac{C_1}{T_2} \left\lceil \frac{T_2}{T_1} \right\rceil, \text{ where } C_1 = T_2 - T_1 \left\lfloor \frac{T_2}{T_1} \right\rfloor$$

$$U = \begin{cases} 1 & \text{if } f = 0 \\ \frac{I + f^2}{I + f} & \text{otherwise} \end{cases}$$

LUB of Processor Utilization (cont'd)

Proof (cont'd):

■ Minimum U ?

$$U = \frac{1+f^2}{1+f}, \text{ where } I=1$$

$$\frac{dU}{df} = \frac{f^2 + 2f - 1}{(1+f)^2} = 0 \Rightarrow \begin{matrix} f_1 = -1 - \sqrt{2} \\ f_2 = -1 + \sqrt{2} \end{matrix} \Rightarrow U = 2(\sqrt{2} - 1), \text{ where } f = -1 + \sqrt{2}$$

About 83 %

LUB of Processor Utilization (cont'd)

Theorem 4.

For a set of n tasks with fixed priority assignment, and the restriction that the ratio between any two request periods is less than 2, the least upper bound to the processor utilization factor is $U=n(2^{1/n}-1)$.

Proof: From the restriction and the proof of theorem 3, we intend to show:

$$C_i = T_{i+1} - T_i, \text{ and } C_n = 2T_1 - T_n \quad (T_n < 2T_1)$$

(1) When $C_1 > T_2 - T_1$


Suppose that $C_1 = T_2 - T_1 + \Delta$, $\Delta > 0$.

Let $C'_1 = T_2 - T_1$, $C'_2 = C_2 + \Delta$, and $C'_i = C_i$ where $3 \leq i \leq n$

C'_1, C'_2, \dots, C'_n also fully utilize the processor.

Let U' denote the corresponding utilization factor,

$$U - U' = \frac{\Delta}{T_1} - \frac{\Delta}{T_2} > 0$$



LUB of Processor Utilization (cont'd)

Proof (Cont'd):

(b) When $C_1 < T_2 - T_1$


Suppose that $C_1 = T_2 - T_1 - \Delta$, $\Delta > 0$.

Let $C_1'' = T_2 - T_1$, $C_2'' = C_2 - 2\Delta$, and $C_i'' = C_i$ where $3 \leq i \leq n$

$C_1'', C_2'', \dots, C_n''$ also fully utilize the processor.

Let U'' denote the corresponding utilization factor,

$$U - U'' = -\frac{\Delta}{T_1} + \frac{2\Delta}{T_2} > 0$$



LUB of Processor Utilization (cont'd)


Proof (Cont'd):

From (1) and (2), we get $C_1 = T_2 - T_1$ if U is the minimum utilization factor.

Similarly, we get

$$\begin{cases} C_2 = T_3 - T_2 \\ \dots \\ C_{n-1} = T_n - T_{n-1} \\ C_n = T_n - 2(C_1 + \dots + C_{n-1}) = 2T_1 - T_n \end{cases}$$

$T_n - T_1$



$$U = \sum_{i=1}^n \frac{C_i}{T_i} = \frac{T_2 - T_1}{T_1} + \dots + \frac{T_n - T_{n-1}}{T_{n-1}} + \frac{2T_1 - T_n}{T_n}$$

Let $R_i = \frac{T_{i+1}}{T_i}$,

$$U = \sum_{i=1}^{n-1} R_i + \frac{2}{R_1 R_2 \dots R_{n-1}} - n \quad (\because R_1 R_2 \dots R_{n-1} = \frac{T_n}{T_1})$$

LUB of Processor Utilization (cont'd)

Proof (Cont'd):

For $U = \sum_{i=1}^{n-1} R_i + \frac{2}{R_1 R_2 \dots R_{n-1}} - n$, to minimize U over $R_k, k = 1, \dots, n-1$

we have:
$$\frac{\partial U}{\partial R_k} = 1 - \frac{2}{R_k R_1 \dots R_{n-1}} = 0$$

Defining $P = R_1 R_2 \dots R_{n-1}$, $R_k P = 2$ for $\forall k, 1 \leq k \leq n-1$

$$\therefore R_1 = R_2 = \dots = R_{n-1} = 2^{1/n}$$

It follows that,

$$U = (n-1)2^{1/n} + \frac{2}{2^{(1-1/n)}} - n = n(2^{1/n} - 1)$$

- For large n , $U \approx \ln 2 = 0.693$

LUB of Processor Utilization (cont'd)

Theorem 5. LUB for a set of n tasks

For a set of n tasks with fixed priority assignment, the least upper bound to the processor utilization factor is $U = n(2^{1/n} - 1)$.

Proof: For details, refer the paper.

If a set of tasks fully utilizes the processor and for some $i, i < m$

$$T_m/T_i \geq 2$$

then, construct another set of tasks that will fully utilize the processor,

$$T_m/T_i < 2$$

- ➔ Show the utilization of the new task that is less than the original one.
- ➔ Hence, we need only consider tasks sets in which the ratio between any two periods is less than 2.

Sample Problem

■ Sample Problem: Applying UB Test

	C	T	U
Task 1	20	100	0.200
Task 2	40	150	0.267
Task 3	100	350	0.286

$U(1) = 1.0$, $U(2) = 0.828$, $U(3) = 0.779$, $U(\infty) = \ln 2 = 0.693$
only sufficient condition

Total utilization:

$$0.200 + 0.267 + 0.286 = 0.753 \\ < U(3) = 0.779$$

The periodic tasks in the sample problem are schedulable according to the UB test.

Exercise

- Exercise: Applying the UB Test Given:

Task	C	T	U
τ_1	1	4	
τ_2	2	6	
τ_3	1	10	

- What is utilization for each task?
- Is the task set schedulable?
- Draw the timeline.
- What is the total utilization if $C_3=2$?

Limitations

- Feasibility test based on utilization bound is sufficient to guarantee the feasibility of any task set, but it is not necessary.
 - UB test has three possible outcomes.
 - $0 \leq U \leq LUB \rightarrow$ success
 - $LUB < U \leq 1.00$ \rightarrow inconclusive
 - $1.00 < U \rightarrow$ overloaded
 - UB test is conservative
 - A more precise test can be applied.
 - E.g) Harmonic task set ($U=1$)

Completion Time Test

Theorem 6. Completion Time Test

For a set of independent, periodic tasks, if each task meets its first deadline, with worst-case task phasing, the deadline will always be met.

→ Completion Time Test

→ Let W_i = completion time of task i . W_i may be computed by the following iterative formula:


$$W_i(n + 1) = C_i + \sum_{j < i} \left\lceil \frac{W_i(n)}{T_j} \right\rceil C_j \text{ where } W_i(0) = 0$$

→ Task i is schedulable if its completion time is before its deadline. That is, $W_i \leq T_i$

Example

- Task τ_1 : $C_1=20$, $T_1=100$, $U_1=0.2$
- Task τ_2 : $C_2=40$, $T_2=150$, $U_2=0.267$
- Task τ_3 : $C_3=100$, $T_3=350$, $U_3=0.286$

- Total utilization: $0.753 \leq U_{\text{lub}}(3) = 3(2^{1/3}-1) = 0.779$
 - 24.7% of the CPU is usable for lower priority background computation.
- Suppose that C_1 is increased to 40,
 - The total utilization is increased to $0.953 > U_{\text{lub}}(3)$.



Example

■ Applying CT Test

- Taking the sample problem, we increase the compute time of τ_1 from 20 to 40. Is the task set still schedulable?
- Utilization of first two tasks: $0.667 < U(2) = 0.828$
 - first two tasks are schedulable by utilization bound test
- Utilization of all three tasks: $0.953 > U(3) = 0.779$
 - utilization bound test is inconclusive
 - need to apply completion time test

Example

- Use CT test to determine if task 3 meets its first deadline

$$W_3(1) = C_3 + \sum_{j<3} \left\lceil \frac{0}{T_j} \right\rceil C_j = C_3 = 100$$

$$\begin{aligned} W_3(2) &= C_3 + \sum_{j<3} \left\lceil \frac{100}{T_j} \right\rceil C_j \\ &= 100 + \left\lceil \frac{100}{100} \right\rceil (20) + \left\lceil \frac{100}{150} \right\rceil (40) = 160 \end{aligned}$$

$$W_3(3) = 100 + \left\lceil \frac{160}{100} \right\rceil (20) + \left\lceil \frac{180}{150} \right\rceil (40) = 220$$

Example

$$W_3(3) = 220$$

$$W_3(4) = 100 + \left\lceil \frac{220}{100} \right\rceil (20) + \left\lceil \frac{260}{150} \right\rceil (40) = 240$$

$$W_3(4) = 100 + \left\lceil \frac{240}{100} \right\rceil (20) + \left\lceil \frac{300}{150} \right\rceil (40) = 240 \Rightarrow \textit{Done!}$$

$$W_3 = 240 < T_3 = 350$$

Task 3 is schedulable using CT test.

Example

■ Exercise: Applying CT Test

Task τ_1 : $C_1 = 1$ $T_1 = 4$

Task τ_2 : $C_2 = 1$ $T_2 = 4$

Task τ_3 : $C_3 = 1$ $T_3 = 4$

- a. Apply UB test.
- b. Draw timeline.
- c. Apply CT test.

Summary

- Summary
 - Utilization bound test is simple but conservative
 - Completion time test is more exact but also more complicated
 - To this point, UB and CT tests share the same limitations.
 - all tasks run on a single processor
 - all tasks periodic and noninteracting
 - deadlines always at the end of the period
 - no interrupts
 - rate monotonic priorities assigned
 - zero context switch overhead
 - tasks do not suspend themselves

Real-Time Systems

J. Lehoczky et al, "The Rate Monotonic Scheduling Algorithm: Exact Characterization And Average Case Behavior", In *Proceedings of the IEEE Real-Time Systems Symposium*, 1989

Motivation

- Utilization bound feasibility test is simple but
 - The feasibility condition is sufficient but not necessary (pessimistic).
 - The average case behavior is substantially better than the worst case behavior.

Problem Formulation

- Task set $\tau = \{ \tau_1, \tau_2, \dots, \tau_n \}$
 - A task is characterized by (T_i, C_i, I_i)
 - T_i : period, C_i : computation time, I_i : Phasing (offset)
 - Each task instance is initiated (released) at times $I_i + kT_i$, $k \geq 0$.
 - The deadline of $(k+1)^{\text{st}}$ instance is $I_i + (k+1)T_i$.
- Basic idea
 - For a set of independent, periodic tasks, if each task meets its first deadline, with the worst-case task phasings, the deadline will always be met.
 - Worst-case task phasing = critical instant ($I_i = 0$ for $1 \leq i \leq n$)

Necessary and Sufficient conditions

- Assume that the task phrasings are all zero (i.e., the first iteration of each task is released at time zero)
- $T_1 < T_2 < T_3 < \dots < T_i$
- τ_1 can be feasibly scheduled when $C_1 \leq T_1$
- If we can find some t in $[0, T_2]$ that satisfy the following condition

$$t \geq \left\lceil \frac{t}{T_1} \right\rceil C_1 + C_2$$

τ_2 can be feasibly scheduled.

- Generally, τ_n can be feasibly scheduled using RM iff we can find some t in $[0, T_n]$ that satisfy the following condition

$$t \geq \sum_{j=1}^n C_j \left\lceil \frac{t}{T_j} \right\rceil$$

Problem Formulation

$$\text{Defining } W_i(t) = \sum_{j=1}^i C_j \left\lceil \frac{t}{T_j} \right\rceil$$

$$L_i(t) = \frac{W_i(t)}{t}, L_i = \min_{\{0 \leq t \leq T_i\}} L_i(t), L = \max_{\{1 \leq i \leq n\}} L_i$$

■ Theorem 1.

- A task τ_i is schedulable using RM if and only if $L_i \leq 1$.
- The entire task set is schedulable using RM if and only if $L \leq 1$.

→ Should we check the condition at every t ?

- The right-hand side of the above equation has jumps only at multiple of T_j
- Practically, we only need to compute $W_i(t)$ at the arrival times of tasks with higher priority than τ_i before the deadline of τ_i and the deadline of τ_j .

Problem Formulation

- $W_i(t)$ is constant, except at a finite number of points when tasks are released. We only need to compute $W_i(t)$ at the scheduling points


Rate Monotonic Scheduling Points

$$S_i = \{kT_j \mid j = 1, \dots, i; k = 1, \dots, \lfloor T_i / T_j \rfloor\}$$

- Theorem 2.
 - A task τ_i is schedulable using RM if and only if

$$L_i = \min_{t \in S_i} \frac{W_i(t)}{t} \leq 1$$

Example : Using Schedulability Points



Example Revisited

- Task τ_1 : $C_1=40$, $T_1=100$, $U_1=0.4$
- Task τ_2 : $C_2=40$, $T_2=150$, $U_2=0.267$
- Task τ_3 : $C_3=100$, $T_3=350$, $U_3=0.286$

- With the schedulability test proposed in the Theorem 2,
 - For task τ_3 , $i=3$ and $S_3=\{100, 150, 200, 300, 350\}$

	$C_1+C_2+C_3 \leq T_1$	$40+40+100 > 100$
or	$2C_1+C_2+C_3 \leq T_2$	$80+40+100 > 150$
or	$2C_1+2C_2+C_3 \leq 2T_1$	$80+80+100 > 200$
or	$3C_1+2C_2+C_3 \leq 2T_2$	$120+80+100 = 300$
or	$4C_1+3C_2+C_3 \leq T_3$	$160+120+100 > 350$

Problem

i	C_i	T_i	i	C_i	T_i
1	20	100	3	80	210
2	30	150	4	100	400

1. L_i ?
2. RM schedulable condition of τ_1 ?
3. RM schedulable condition of τ_2 ?
4. RM schedulable condition of τ_3 ?
5. RM schedulable condition of τ_4 ?

More Results on RM

- A stochastic analysis for a randomly generated set of periodic tasks scheduled by RM has shown that
 - The average scheduling bound is usually much better than worst case behavior.

Reading List

- Least Slack Time Rate First: an Efficient Scheduling Algorithm for Pervasive Computing Environment